

- **FastAPI** → A **Python web framework** to build **HTTP APIs** (REST/JSON) for humans, services, browsers, etc. It's built on Starlette/Pydantic with automatic OpenAPI docs. [\[fastapi.tiangolo.com\]](https://fastapi.tiangolo.com), [\[fastapi.tiangolo.com\]](https://fastapi.tiangolo.com)
  - **FastMCP** → A **Python framework for Model Context Protocol (MCP)** to build **servers and clients that expose tools/resources/prompts to LLM hosts** (e.g., Claude Desktop, IDE assistants). It focuses on LLM-to-tool integration and can even generate OpenAPI/FastAPI surfaces from MCP tools. [\[fastmcp.wiki\]](https://fastmcp.wiki), [\[pypi.org\]](https://pypi.org)
- 

## Core intent & audience

Aspect	FastAPI	FastMCP
Primary goal	Build high-performance <b>HTTP APIs</b>	Build <b>MCP servers/clients</b> that let <b>LLMs</b> call your tools and read resources
Main consumers	Browsers, microservices, mobile apps, human developers	<b>LLM hosts/clients</b> (Claude Desktop, editors, agent runtimes)
Typical output	JSON over HTTP with OpenAPI docs	MCP primitives: <b>Tools</b> (actions), <b>Resources</b> (read-only data), <b>Prompts</b> (templates) exposed via JSON-RPC over stdio/SSE transports

- FastAPI's charter is conventional API development with standards like **OpenAPI/JSON Schema**, interactive Swagger/ReDoc docs, and dependency injection for request handling. [\[fastapi.tiangolo.com\]](https://fastapi.tiangolo.com), [\[fastapi.tiangolo.com\]](https://fastapi.tiangolo.com)
  - FastMCP's charter is **LLM integration** using the MCP standard (JSON-RPC 2.0, stdio and SSE transports) so models can **discover** and **invoke** your capabilities safely. [\[modelcon...rotocol.io\]](https://modelcon...rotocol.io), [\[en.wikipedia.org\]](https://en.wikipedia.org)
- 

## Architecture & transport

- **FastAPI** runs an **ASGI** app (often with Uvicorn) serving HTTP endpoints. Everything is request/response over HTTP. [\[devdocs.io\]](https://devdocs.io)
  - **FastMCP** implements MCP's client/server protocol. Communication is **JSON-RPC** over **stdio** (local) or **SSE/HTTP** (remote). It manages capability discovery, resource listing, and tool execution semantics tailored for LLMs. [\[modelcon...rotocol.io\]](https://modelcon...rotocol.io), [\[apxml.com\]](https://apxml.com)
- 

## Primitives vs. endpoints

- **FastAPI**: You define `@app.get`, `@app.post` routes; models validated by Pydantic; OpenAPI is auto-generated. [\[fastapi.tiangolo.com\]](https://fastapi.tiangolo.com)

- **FastMCP**: You decorate Python functions with `@mcp.tool` (actions), expose **resources** (read-only context, like GET), and **prompts** for reusable interaction templates—all discoverable by the LLM host. [\[fastmcp.wiki\]](#)
- 

## Documentation & generation

- **FastAPI** gives you **interactive docs** (Swagger UI/ReDoc) out of the box for your HTTP API. [\[fastapi.tiangolo.com\]](#)
  - **FastMCP** provides MCP docs and, notably, **can generate OpenAPI/FastAPI layers** from MCP definitions when needed—handy if you want both LLM-facing MCP **and** human-facing REST. [\[pypi.org\]](#)
- 

## Security & auth focus

- **FastAPI**: Built-in patterns for OAuth2/JWT, dependency injection for auth, typical web API security concerns. [\[geeksforgeeks.org\]](#)
  - **FastMCP**: Emphasizes **enterprise authentication** (Google, GitHub, Azure, Auth0, etc.) for MCP apps and **user-governed tool execution** (approval flows) typical in LLM environments. [\[pypi.org\]](#), [\[modelconte...rotocol.be\]](#)
- 

## Typical use cases

### Choose FastAPI when you need:

- A REST/JSON API for your web/mobile app or microservices.
- Automatic docs, validation, and high performance for conventional HTTP workloads. [\[fastapi.tiangolo.com\]](#)

### Choose FastMCP when you need:

- An **LLM-aware interface** so assistants can call your functions (tools), read your data (resources), and use prompts—e.g., an AI DevOps agent that triggers pipelines or reads logs. [\[fastmcp.wiki\]](#), [\[datacamp.com\]](#)
- 

## How they can complement each other

### You can **combine** them:

- Implement your AI-facing layer in **FastMCP** (expose tools/resources to LLMs).

- Expose parts of the same functionality to humans or services via **FastAPI** (or generate an OpenAPI/FastAPI surface from FastMCP). [\[pypi.org\]](https://pypi.org)
- 

### Quick decision guide

- Are your primary callers humans/services over HTTP? → **FastAPI**.  
[\[fastapi.tiangolo.com\]](https://fastapi.tiangolo.com)
- Are your primary callers LLM hosts/assistants needing tool/resource discovery & safe invocation? → **FastMCP**. [\[modelconventionrotocol.io\]](https://modelconventionrotocol.io)
- Need both? Start with **FastMCP**, then generate an HTTP surface (or handcraft with **FastAPI**). [\[pypi.org\]](https://pypi.org)